


With Pre-paging

Pre-paging		
Page Size	Algorithm	Page Swaps
1	FIFO	102416
2	FIFO	100214
4	FIFO	106486
8	FIFO	120172
16	FIFO	177156
1	LRU	99702
2	LRU	94318
4	LRU	93190
8	LRU	92505
16	LRU	103298
1	CLOCK	99817
2	CLOCK	94393
4	CLOCK	93599
8	CLOCK	96758
16	CLOCK	177151

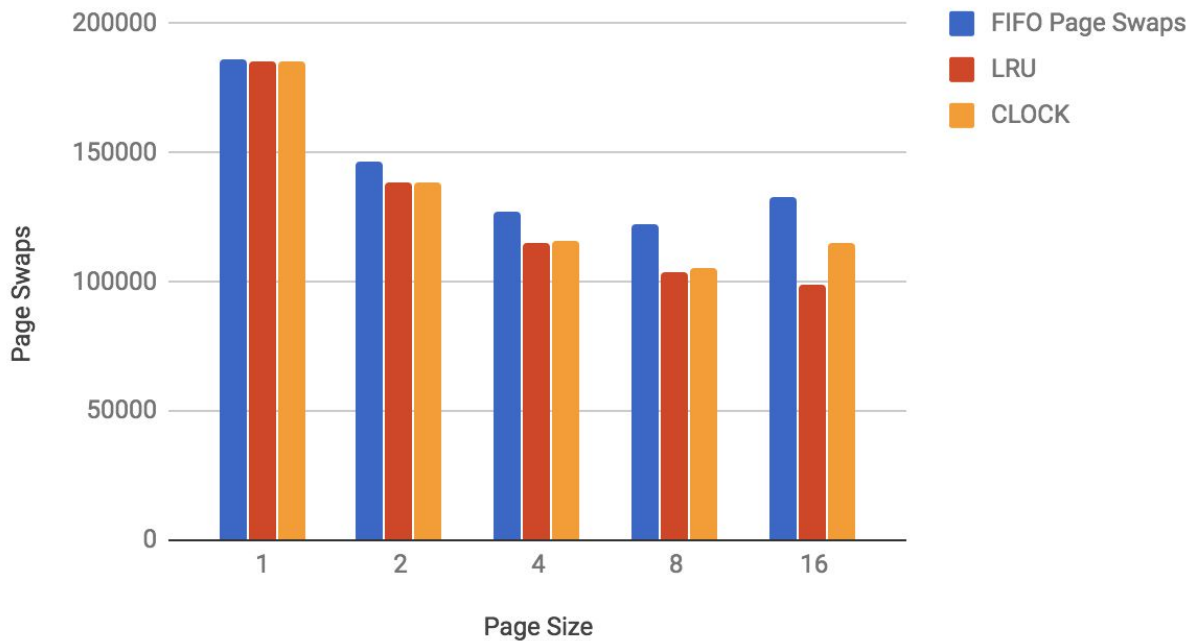
Pre-paging



With Demand Paging

Demand paging		
Page Size	Algorithm	Page Swaps
1	FIFO	186384
2	FIFO	146739
4	FIFO	127322
8	FIFO	122718
16	FIFO	133157
1	LRU	185063
2	LRU	138449
4	LRU	115484
8	LRU	103543
16	LRU	98930
1	CLOCK	185098
2	CLOCK	138501
4	CLOCK	115730
8	CLOCK	105376
16	CLOCK	115238

Demand Paging



With the number of page swaps as a metric, LRU is the most efficient algorithm for virtual page management, with the CLOCK algorithm being a close second and FIFO the least.

LRU imposes significant overhead and is difficult to implement, but is based on the principle of locality, in that the page selected for replacement is the one that has not been referenced for the longest time.

The **CLOCK** algorithm takes advantage of a circular implementation for finding pages to replace. This implementation has less overhead, but could take longer to find a page to replace, and is hence beat out by LRU closely.

The **FIFO** algorithm had the simplest implementation by far, and required very little overhead. It replaces the page that had been in memory the longest every time. The greater the amount of page frames, the more page swaps it has to do and hence the more poorly it performs.

With **demand paging**, pages are brought into memory only when a location on the page is referenced during execution. With **pre-paging**, other pages are also brought into memory based on common access patterns, locality, or time last referenced with the hope that it will be needed in the near future and will hence not require another access to main memory.

Prepaging was more efficient than demand paging for page-sizes less than 8. They performed similarly for a page size of 8, and then demand paging took the lead for page sizes of 16 and up. This behavior matches our expectations given the descriptions of the differences in the algorithms above.

Random memory access traces would have significantly reduced performance, as the principle of locality states that, with pre-paging, we can reduce the number of memory accesses by taking advantage of memory locations being adjacent to one another. Using larger page sizes would not help to reduce page faults, as demonstrated by the above graphs.